

General Q&A

No questions.

Theory Tasks

Chapter 6: Association Rules

Why can support/confidence for association rules be misleading?

- support:
 - measure that can also be applied to plain itemsets, not only association rules
 - quantifies how often itemsets on LHS and RHS of association rule occur together
 - i.e., tells how often rule satisfied (relative to number of transactions or absolute)
 - but does not consider how often rule is invalid (only LHS satisfied) → use confidence
- confidence:
 - normalizes support of association rule with support of LHS
 - conceptually similar to metric *precision* in classification
 - asymmetric measure: does not consider support of RHS
 - * might be high even if there is negative correlation between LHS and RHS
 - * in particular, danger if support of RHS is high already without LHS
 - * see basketball/cereal example from lecture
 - in contrast, *lift* is symmetric quality measure (but has other weaknesses)

Do you see a relationship between association rules and decision trees?

- at first glance, very different approaches
 - association rules are unsupervised (no target variable)
 - * decision trees are supervised
 - association rules work with transactions of (feature-less and target-less) items
 - * decision trees work with feature-target representation of data objects
 - association rules might not apply (if LHS of rule not contained in transaction)
 - * decision trees classify each data object (they partition data space fully)
 - association rules are individual rules
 - * decision trees are a hierarchical structure of rules (applied in combination)
- main similarity: both approaches use if-then rules
 - association rules: if transaction contains LHS, then it also should contain RHS
 - decision tree: if feature has certain value, follow corresponding path and either test another feature (internal node) or assign class label (leaf node)
- you can apply association rules to tabular feature-target data
 - step 1: transform dataset into transaction-item representation
 - * each data object becomes a transaction
 - * each feature value or target value becomes an item
 - * thus, each transaction has same length
 - * don't forget feature/target name → multi-dimensional itemset (see the lecture)
 - * need to discretize features (unlike decision trees, which support numeric features)
 - * exemplary transaction for a data object from the *iris* dataset: {sepal.length=short, sepal.width=long, petal.length=short, petal.width=short, species=setosa}
 - * i.e., values of four features and a target transformed to transaction with five items

- step 2: mine multi-dimensional association rules
 - * rules like {petal.width=short} → {petal.width=long} cannot occur
 - * also, we want rules where feature values are on LHS and a target value is on RHS
 - * (we could also mine rules to explore relationships between features)
 - * these constraints might make mining (candidate generation) more efficient
- step 3: post-process rules
 - * problem: zero, one, or multiple rules might apply to features of a data object
 - * need to define default prediction if no rule matches
 - * need to decide what to do if multiple rules match, e.g., bring rules into an order
- exemplary algorithm: CBA (Classification Based on Associations) from paper “[Integrating Classification and Association Rule Mining](#)”
- exemplary Python package: [pyarc](#)

How does sorting help in FP-trees?

- sorting by frequency is integral component of original FP-tree algorithm
 - done between first scan of database (count items) and second scan (build tree)
- *prefix path property* of algorithm depends on frequency sorting of items in transactions
 - bonus task: does algorithm work with arbitrary (potentially fixed) order of items?
- sorting by frequency improves efficiency
 - FP-tree smaller (has less paths) if transactions share first few items (so-called prefix)
 - more frequent items more likely to be shared, thus sorted to beginning of transactions