

General Q&A

No questions.

Theory Tasks

Chapter 2: Fundamentals

For data aggregation, is there a relationship between distributive/algebraic/holistic aggregates and self-maintainable aggregates?

- self-maintainability is a property regarding change operations (like insertion) of data
 - basic definition: update aggregate after change only from change and previous value
 - aggregate might be self-maintainable regarding some operations, but not others
- distributive aggregates:
 - can be computed from aggregates of same type determined on data partitions
 - thus, self-maintainable regarding insertion
 - * compute aggregate on set of new point(s)
 - * combine with value of aggregate on existing points
 - * e.g., min, max, count, and sum can be updated with new point(s)
 - sum / count also self-maintainable regarding deletion
 - * subtract values of deleted points / subtract number of deleted points
 - min / max not self-maintainable regarding deletion
 - * no problem if deleted value higher than min / lower than max → aggregate same
 - * problem if deleted value equal to min / max → new value of aggregate unclear
- algebraic aggregates:
 - can be computed from set of distributive aggregates
 - self-maintainability in narrow sense violated (need multiple sub-aggregates)
 - self-maintainable in broad sense (if access to all necessary sub-aggregates)
 - e.g., if sum and number of points stored, can still compute mean after insert/delete
- holistic aggregates:
 - cannot be computed from fixed amount of sub-aggregates (e.g., median, mode)
 - thus, not self-maintainable regarding insertion/deletion

Chapter 7: Clustering

What are ‘convex space partitions’ in the context of k -Means?

- all clustering algorithms typically partition space (at least implicitly)
 - for each point, you can decide to which cluster it belongs
 - for ‘soft’ clustering methods like mixture models, membership is probabilistic
 - for k -Means, points are assigned to closest centroid
 - partitioning can be visualized as [Voronoi diagram](#)
- partitions in k -Means are convex, as
 - all points on a hypersphere around centroids have same distance to centroids
 - hyperspheres are convex
 - k -Means attempts to minimize distances to centroids

How can we improve upon the complexity of k -Medoid?

- complexity according to lecture: $O(t \cdot k \cdot (n - k)^2)$
- can reduce number of iterations t , but that might just prevent proper convergence
- main problem: in each iteration,
 - consider each pair of medoid and non-medoid $\rightarrow O(k \cdot (n - k))$
 - compute change in objective function $\rightarrow O(n - k)$
- improvement 1: CLARA
 - naming: ‘Clustering LARge Applications’
 - run PAM on sample of size $n' < n$
 - replaces n with n' in complexity estimate
- improvement 2: CLARANS
 - naming: ‘Clustering LARge Applications based upon RANdomized Search’
 - use full dataset (so not an extension of CLARA!)
 - randomly select medoid $numlocal$ times (instead $t \cdot k$)
 - for medoid, randomly select non-medoid at most $maxneighbor$ times (instead $n - k$)
 - overall complexity: $O(numlocal \cdot maxneighbor \cdot (n - k))$
- improvement 3 (not in lecture): [FasterPAM](#)
 - same results as PAM, but $O(k)$ speedup due to improved medoid swapping

How does OPTICS work?

- density-based clustering algorithm
- naming: ‘Ordering Points To Identify the Clustering Structure’
- usually for numeric data, but actually combinable with arbitrary dissimilarity measure
- hyperparameters:
 - $MinPts$: min number of points in neighborhood to make point *core* (as in DBSCAN)
 - ϵ : max neighborhood radius (different to DBSCAN, despite same name)
- output: no clustering, but ordered list of points with core/reachability distances
 - can be used to create reachability plot
 - can be used to create fixed, DBSCAN-like clustering (by choosing some $\epsilon^* < \epsilon$)
 - in particular, can extract clusterings with different density thresholds
 - in particular, can use reachability plot to select a sensible ϵ^*
- important definitions:
 - core distance: min neighborhood radius of a point to make it *core*
 - reachability distance: min distance such point directly density reachable from other
 - * at least actual distance between the two points
 - * at least core distance of second point
- procedure:
 - process each point exactly once
 - * retrieve neighbors and compute core distance
 - * write point to output (it is processed now)
 - * check if point *core*; if yes, update reachability distances of unprocessed neighbors
 - maintain priority list
 - * points ordered by minimum reachability distance to all points in output
 - * process points in this order (other than DBSCAN, which uses arbitrary order)
- complexity: $O(n \cdot \log n)$ with spatial index structure for neighborhood queries
 - also, ϵ must not be too high for efficient neighborhood queries
 - $O(n^2)$ with naive implementation (neighborhood queries in $O(n)$)

How does k -Mode work?

- partitioning-clustering algorithm for categorical data
- hyperparameter: number of clusters k
- output: cluster centroids, cluster assignments (latter can be computed with former)
- complexity: $O(t \cdot k \cdot n)$ with t iterations and n points
- procedure similar to k -Means, i.e., repeatedly
 - assign points to their closest centroids
 - re-compute centroids from points assigned to them
- adaptation 1: Hamming distance instead of Euclidean distance
 - reason: data is categorical instead of numeric
 - thus, cannot say how different points are, only if attribute values are different or not
 - Hamming distance sums (binary) differences over all attributes
- adaptation 2: centroid computation
 - for categorical data, cannot compute mean
 - instead, take the mode (most frequent value) of each attribute separately
 - centroid might not be an actual point from dataset (but applies to k -Means as well)

Chapter 8: Outlier Detection

Why should outlier detection be considered unsupervised?

- outlier-detection problems (datasets) typically don't have class labels
 - by definition, that's unsupervised learning
 - in literature, also outlier detection with partially or fully labeled datasets
- outliers don't form a clearly defined target
 - are rare
 - are diverse: usually not only distinct from 'normal' data, but also other outliers
 - i.e., no unique concept/distribution behind them, so characterization difficult

For k -distance-based outliers, what are the dangers of setting k very low or very high?

- k very low: doesn't recognize outliers if close to a few other outliers
 - e.g., for $k = 1$, two outliers close to each other won't be recognized
- k very high: small clusters in data (distant from remaining data) recognized as outliers
 - e.g., imagine $k = 50$ and a cluster of 40 points well-separated from remaining points
- in both examples, depends on use case if you want to label mentioned points as outliers

How do Restricted Boltzmann Machines work for outlier detection?

- two-layer neural networks
 - 'visible' and 'hidden' neurons connected to each other, but not among themselves
- try to learn probability distribution of training data
 - various hyperparameters control training, e.g., number of iterations
 - model (parameters): weight matrix between neurons, biases of neurons
- to score a point as outlier, put it in visible layer and compute 'free energy'
 - says how unlikely point is, given the learned probability distribution (weights)
- compared to other neural-network approaches (Autoencoders, Self-Organizing Maps):
 - architecture different
 - scoring idea (difference to some learned representation of data) roughly similar